



VII

Making Node.js Fast

Optimization and Data Structures

Mattias Hansson, 2022

Story: Cargo Cults



Story: Cargo Cults



Lesson

Don't do Cargo Cult Programming!





VII


Making Node.js Fast

Optimization and Data Structures

Mattias Hansson, 2022


What can we do to speed up a program?

Increase hardware performance



Change programming language


Rewrite the program to do the same with less work
or in a better way



Optimization



Way of Thinking



Practice

Way of Thinking



VS




‘Recipe’



Don't optimize!

“Premature optimization”:

- During PoC, before v 1.0
- The ‘skill trap’
- Not if it's not needed. Ever!





When do we optimize?

Only when it's the best thing we can do right now.

- Usability
- Infrastructure
- As a hobby



Measure, measure, measure!

Always measure (profile) first!





Way of thinking

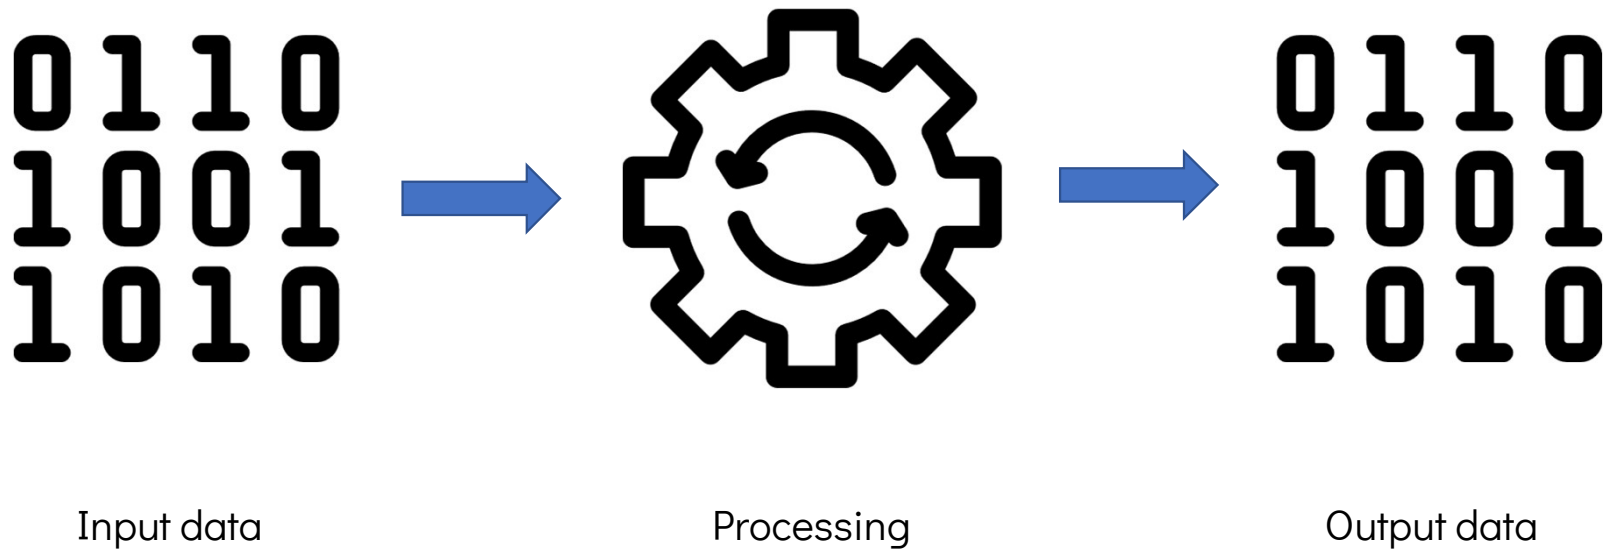
‘What am I asking the computer to do?’

- With the program
- In this function
- In this statement (line)







“The diagram of every program”




Photoshop

- Example: Resize image





Online banking

- Example: Bank statement





ZIP

- Example: compress files



Concept: Data streaming






Speed = total time

0110
1001
1010

Input data



Processing

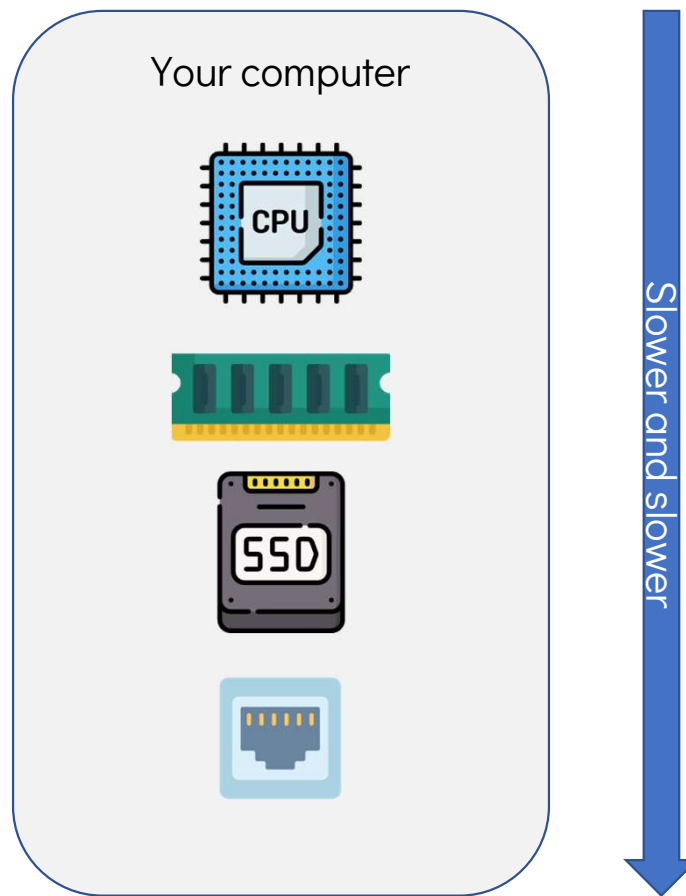


0110
1001
1010

Output data





Computers and Networks





Computers and Networks


Your computer



Your network



The internet



Slower and slower



Latency in computers / networks

Latency

“The amount of waiting time for something to finish.”

Usually with computers:



The seek time “the time to ask for and locate some data”

+



The transfer time “the time it takes to transfer the data”



Latency: Ping





Latency: Download





Latency Numbers Every Programmer Should Know

■ 1 ns

■ L1 cache reference: 0.5 ns

■ Branch mispredict: 5 ns

■ L2 cache reference: 7 ns

■ Mutex lock/unlock: 25 ns

■ = 100 ns

■ Main memory reference: 100 ns

■ = 1 μ s

■ Compress 1 KB with Zippy: 3 μ s

■ = 10 μ s

■ Send 1 KB over 1 Gbps network: 10 μ s

■ SSD random read (1 Gb/s SSD): 150 μ s

■ Read 1 MB sequentially from memory: 250 μ s

■ Round trip in same datacenter: 500 μ s

■ = 1 ms

■ Read 1 MB sequentially from SSD: 1 ms

■ Disk seek: 10 ms

■ Read 1 MB sequentially from disk: 20 ms


■ Packet roundtrip CA to Netherlands: 150 ms

Source: <https://gist.github.com/2841832>




Human understandable latencies

L1 cache reference 0.5 ns **One heart beat (0.5 s)**
L2 cache reference 7 s **Long yawn**



Main memory reference 100 s **Brushing your teeth**



Read 1 MB sequentially from memory 3 days **A long weekend**
SSD random read 1 day, 17 hours **A normal weekend**



Round trip within same datacenter 6 days (local network) **A medium vacation**



Read 1 MB sequentially from SSD 12 days **Waiting for almost 2 weeks for a delivery**

Send packet CA->Netherlands->CA 4 years 10 months **Average time it takes to complete a bachelor's degree**



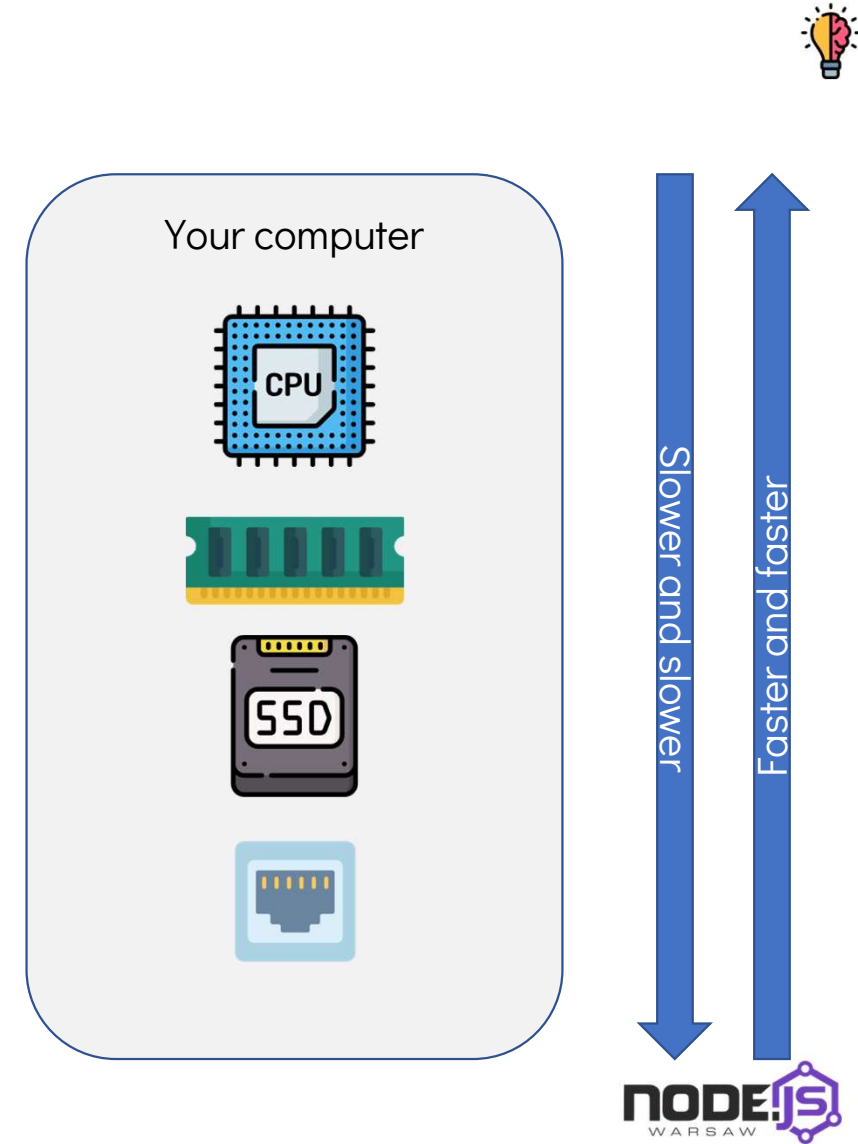
Credit: <https://stackoverflow.com/questions/44485470/understanding-o1-vs-on-time-complexity-intuitively>

Concept #1

Can we bring data closer to the processing?

Example: By using (built in) cache

Demo 1






Concept #2 (algorithm primer)

Can we access the data smarter?

Example: look for value in list






Concept #2 (algorithm primer)

Can we access the data smarter?

Example: binary search


Example: look for 4





Concept #2

Example: look for 24



Not Found!

Demo 2



Way of thinking

‘What am I asking the computer to do?’



```
r = binary_search(arr, search_value) !== -1;
```



```
r = arr.indexOf(search_value) !== -1;
```



Concept #3

Latency multiplies



1000 tasks each taking 1 ms is one second!


1000 simultaneous requests, avg. response time =
1/2 sec!



JavaScript Data Structures

- Lists []
- Objects {}

Lists and Objects work with JSON





JavaScript List [] (aka Array)

- Zero or more items
- Ordered
- Index starts from 0 (zero)
- Access items by index fast
- Items can be other Lists and Objects





JavaScript List (aka Array)

List examples:

```
[]
```

```
[1, 2, 3]
```







```
["aa", "bb", "cc"]
```

```
["a", "b", "c", 342, 9]
```

```
["a", ["hello", "world"], 15, "z"]
```




JavaScript List Characteristics

- Search for item 
- Add item 
 - Add item (push) to the end of the list 
- Delete item 
 - Delete item (pull) from the end of the list 
- Access item 



JavaScript Object {}

- Zero or more properties (Key, Value)
- Keys can be number or string
- Keys are unique
- Values can be primitives (number, string etc.) or other Lists and Objects
- Access by key fast





JavaScript Object



Object examples

```
{ }
```

```
{  
  "name": "Mattias",  
  "height": 189  
}
```



JavaScript Object



Object examples

```
{  
  "bird": ["owl", "hawk"],  
  "bear": ["grizzly", "polar", "grylls"],  
  "dog": ["asian shepherd"]  
}
```




JavaScript Object Characteristics

- Search for item 
- Add item 
- Delete item 
- Access item 



Practice: Airport database



Open data from: <https://ourairports.com/>


Unoptimized

- Lookup the result (by iteration)



Memoization

- Lookup the result
- Store the calculated result for future lookups



(Lazy) Precalculation

- On first lookup, build the precalculated data
- On subsequent lookups, returns precalculated value



Summary



‘What am I asking the computer to do?’



Measure, measure, measure



Learn more about algorithms and data structures



The End

Making Node.js Fast

Optimization and Data Structures

Mattias Hansson, Node.js Warsaw VII, 2022-03-16